



Machine Learning for Embedded Programs Optimisation

Learning Compilers for Configurable Processors

Hugh Leather
University of Edinburgh

Overview

- Members of the Milepost group
- Introduction to machine learning in compilers
- Learning compiler flags for an embedded processor
- Learning compilers over multiple architectures
- List of open source software in the project
- Pros/cons of open source software

Milepost Members



Funded by:
Information Society Technologies (IST) of the
European Union under 6th FWP



Machine Learning in Compilers

- Writing compiler heuristics is hard
 - Need deep knowledge
 - Change elsewhere needs change everywhere
 - Architectures/libraries/OS change
- Humans can't keep up
 - Heuristics remain unoptimised
- Can machine learning help?

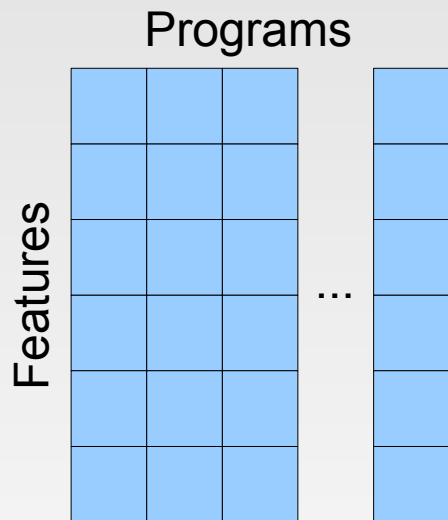
Machine Learning in Compilers

- We start by choosing some information that might be useful to the heuristic (*features*)

number of instructions	
mean dependency depth	
branch count	
loop nest level	
...	
trip count	

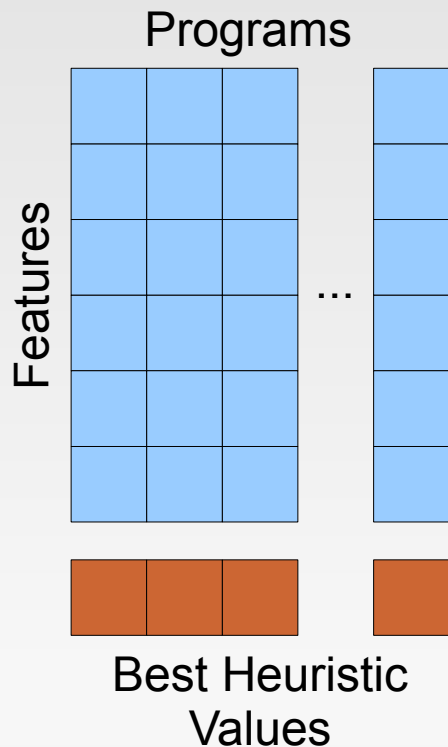
Machine Learning in Compilers

- Collect lot of example benchmarks and compute their features



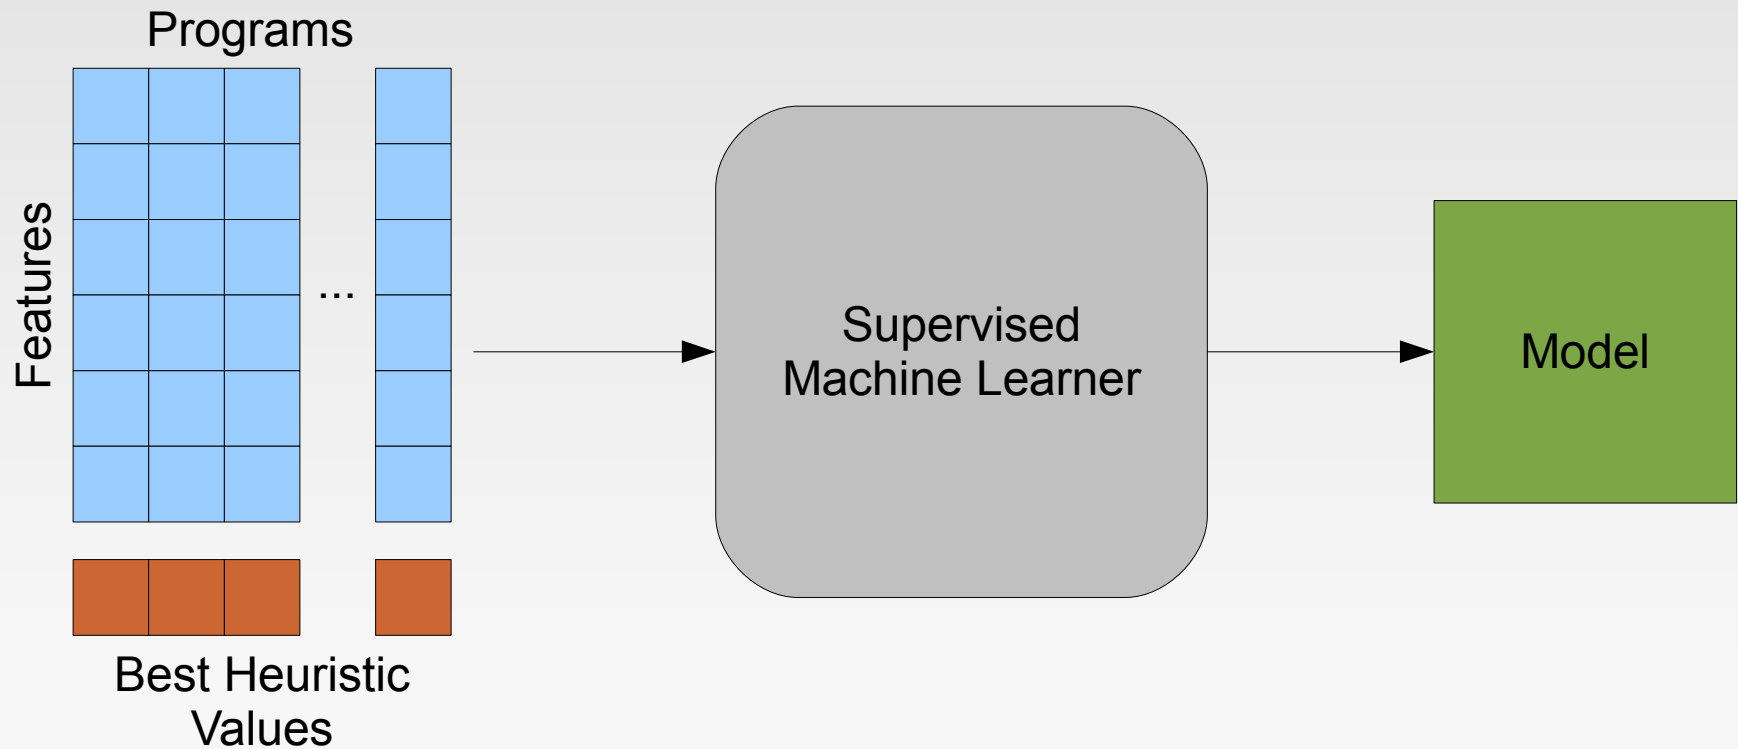
Machine Learning in Compilers

- Execute the programs with different compilation strategies and determine best for each



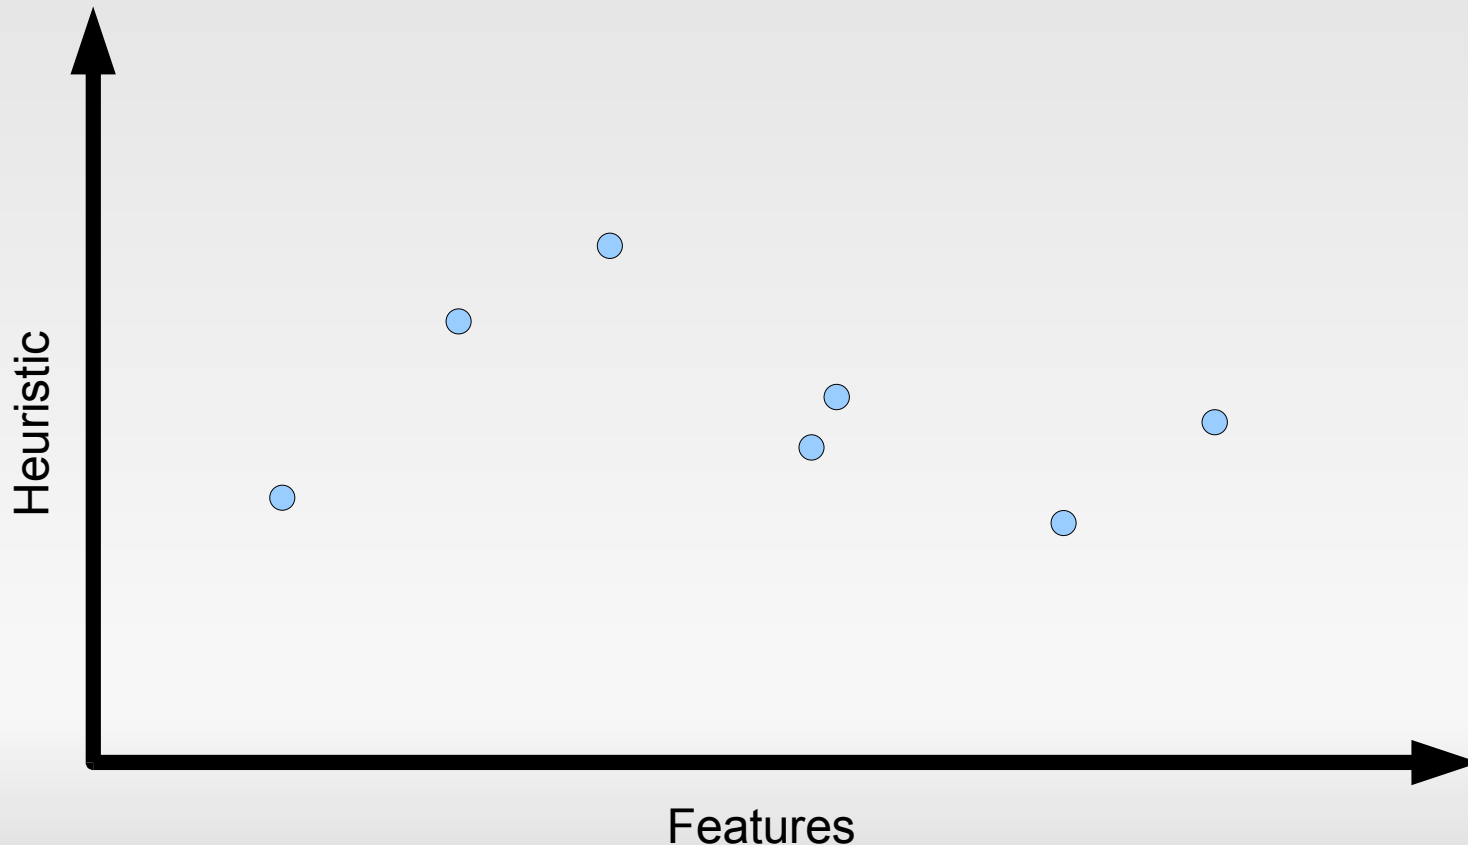
Machine Learning in Compilers

- Pass these to a machine learning tool
- It learns a model



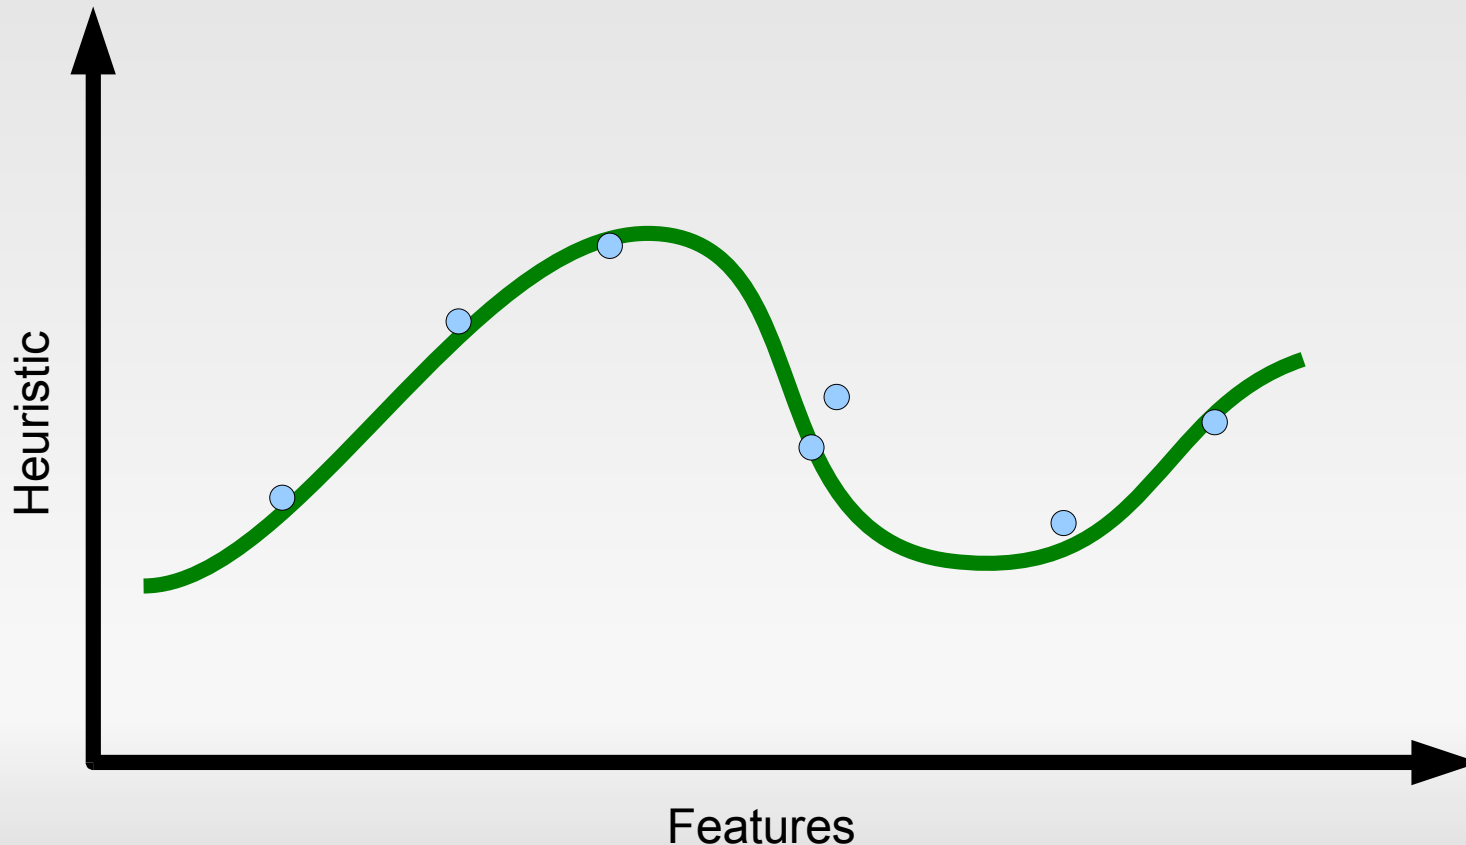
Machine Learning in Compilers

- A model is really just a way of fitting a curve to data



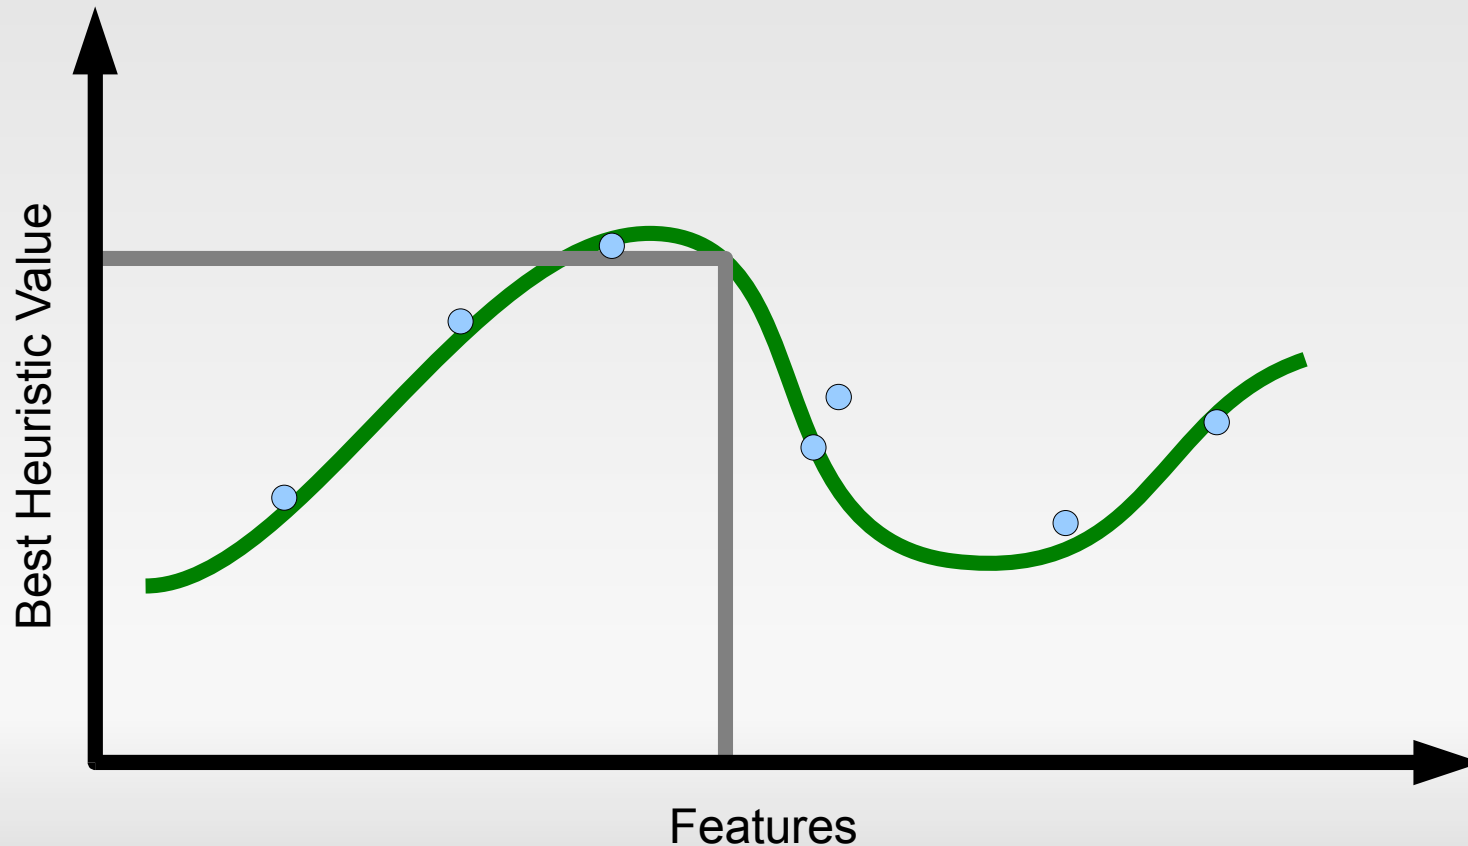
Machine Learning in Compilers

- A model is really just a way of fitting a curve to data



Machine Learning in Compilers

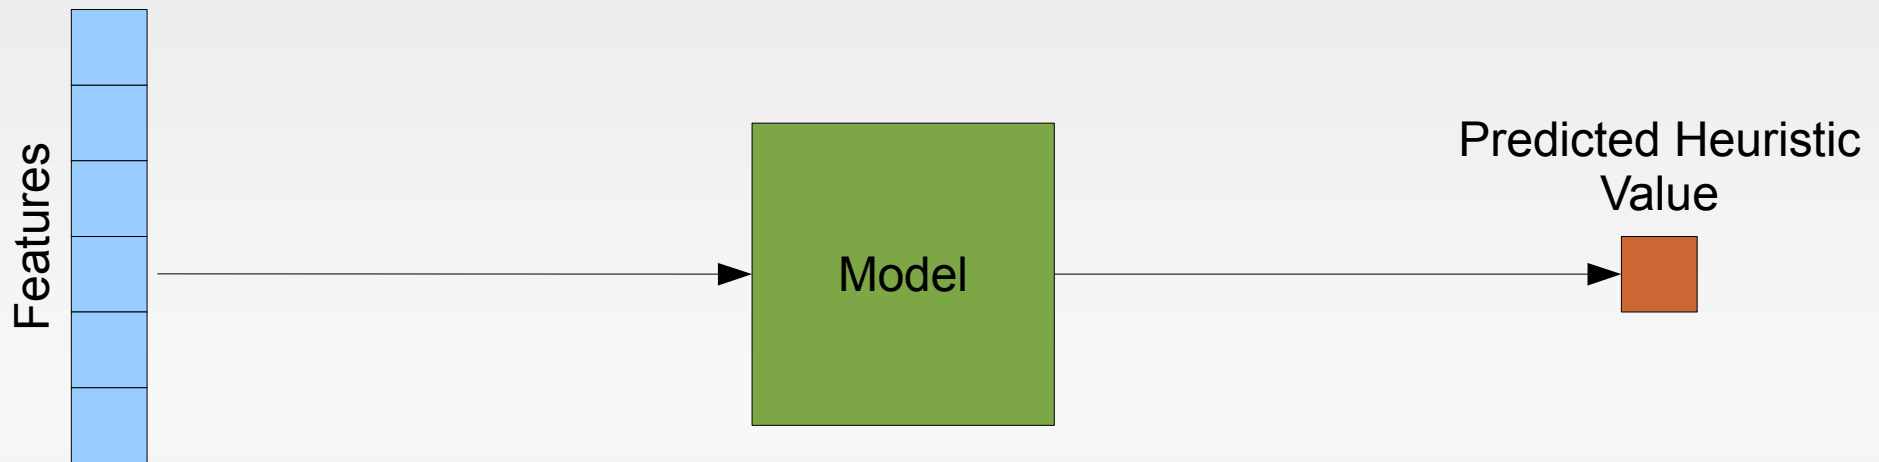
- Gives heuristic for unseen points



Machine Learning in Compilers

- Insert model into compiler – heuristic replaced
- Automatic
 - No human expert required
 - Can be redone after changes

New Program

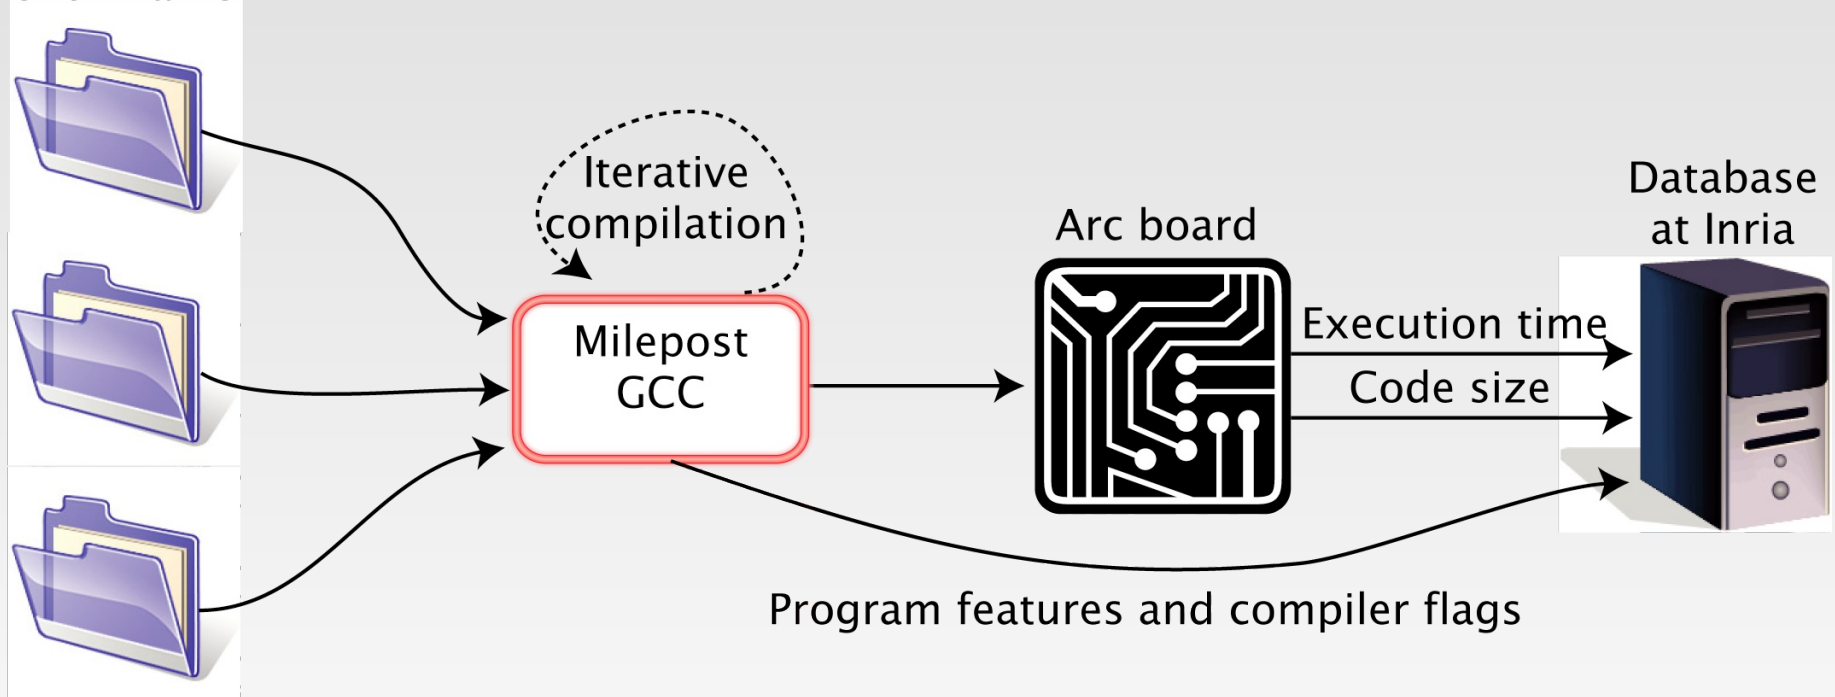


An Experiment

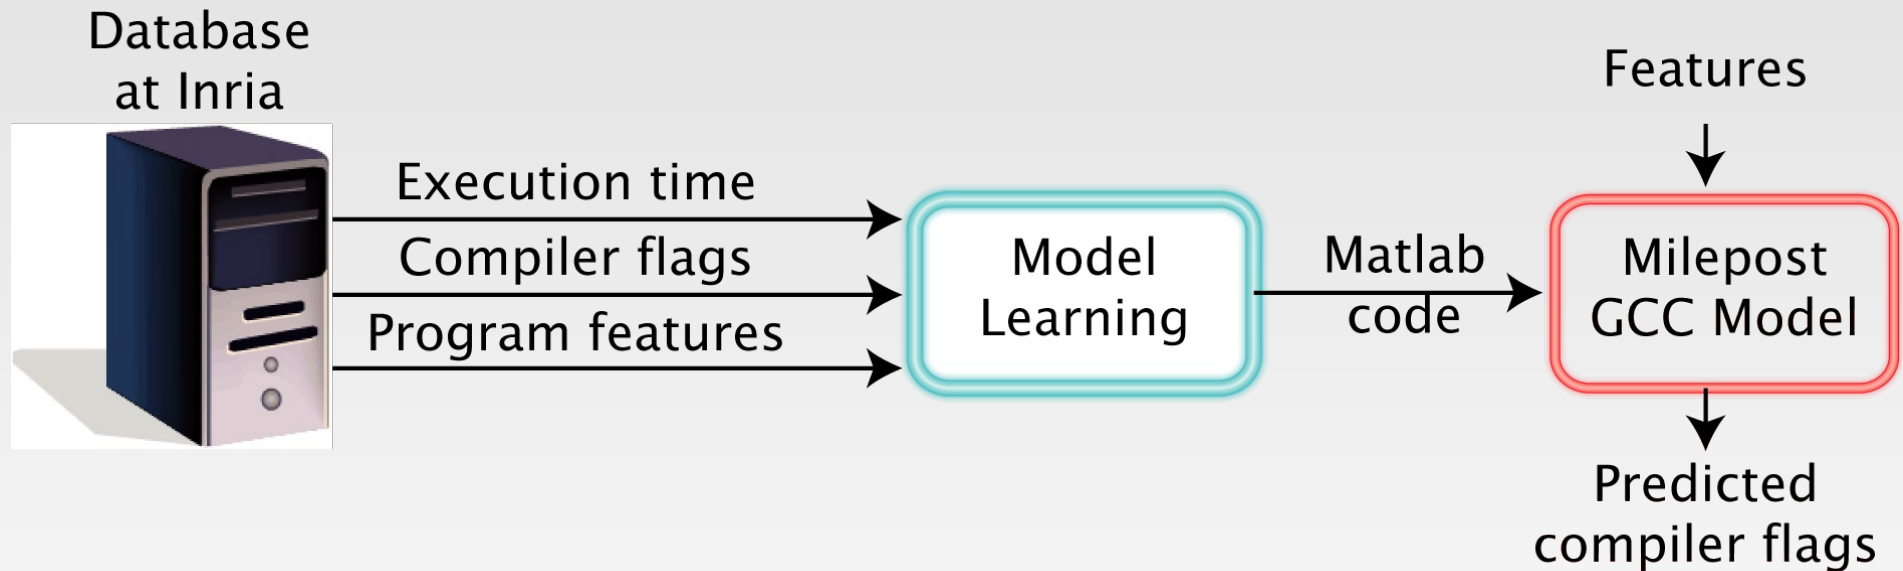
- Learn optimisation flags for MiBench with GCC
- Target the ARC 725 customisable processor
- Trained on
 - 500 random flag settings
 - 5 executions each
 - (2 months runtime on 2 machines)

An Experiment Filling the Database

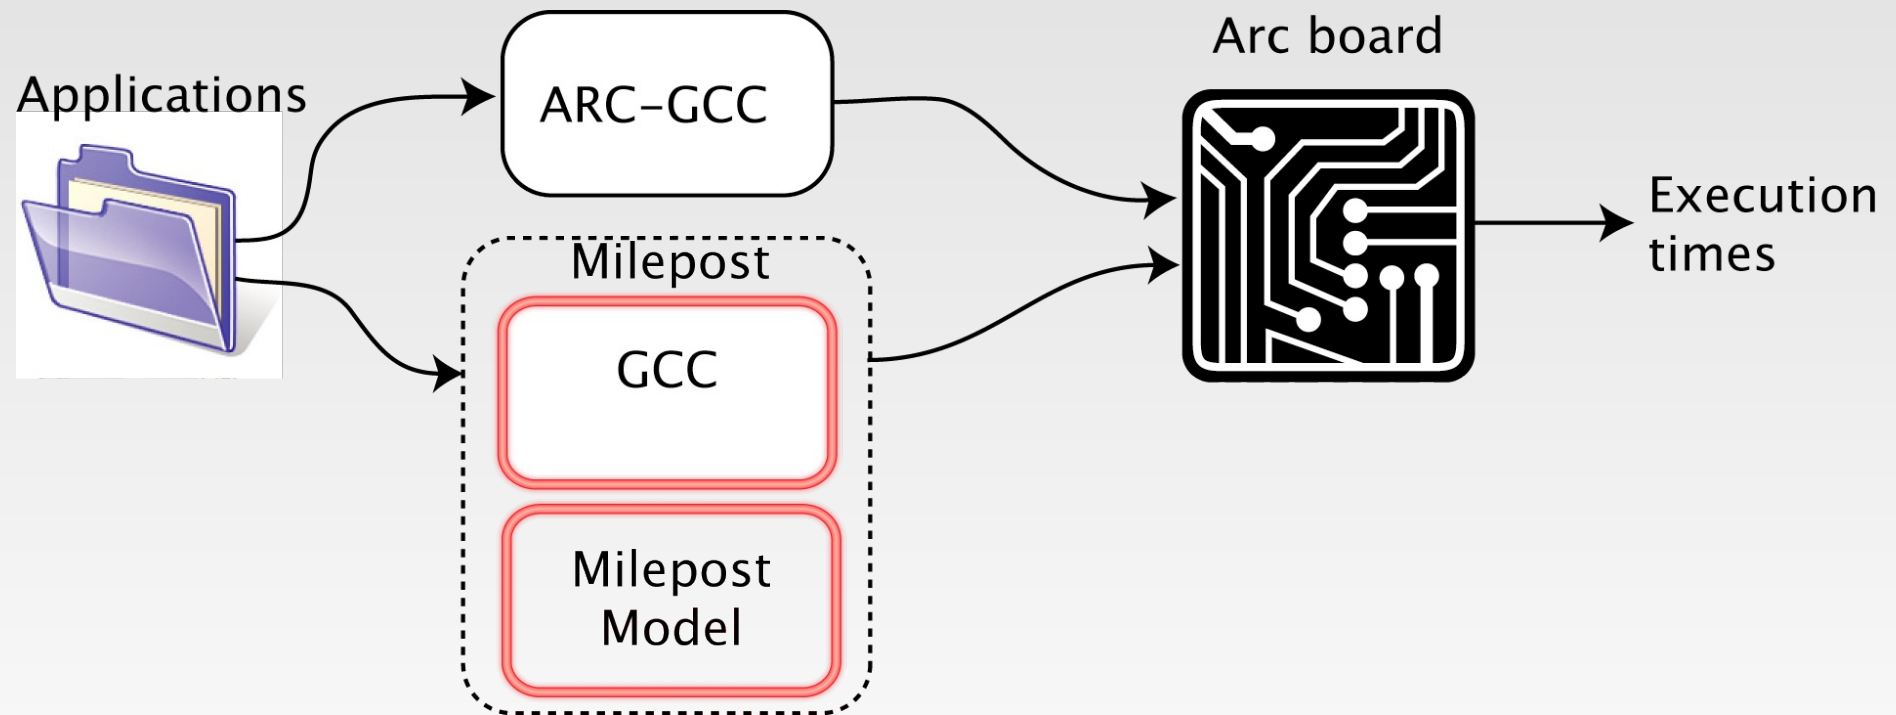
Benchmarks



An Experiment Building the Model



An Experiment Using Improved Compiler



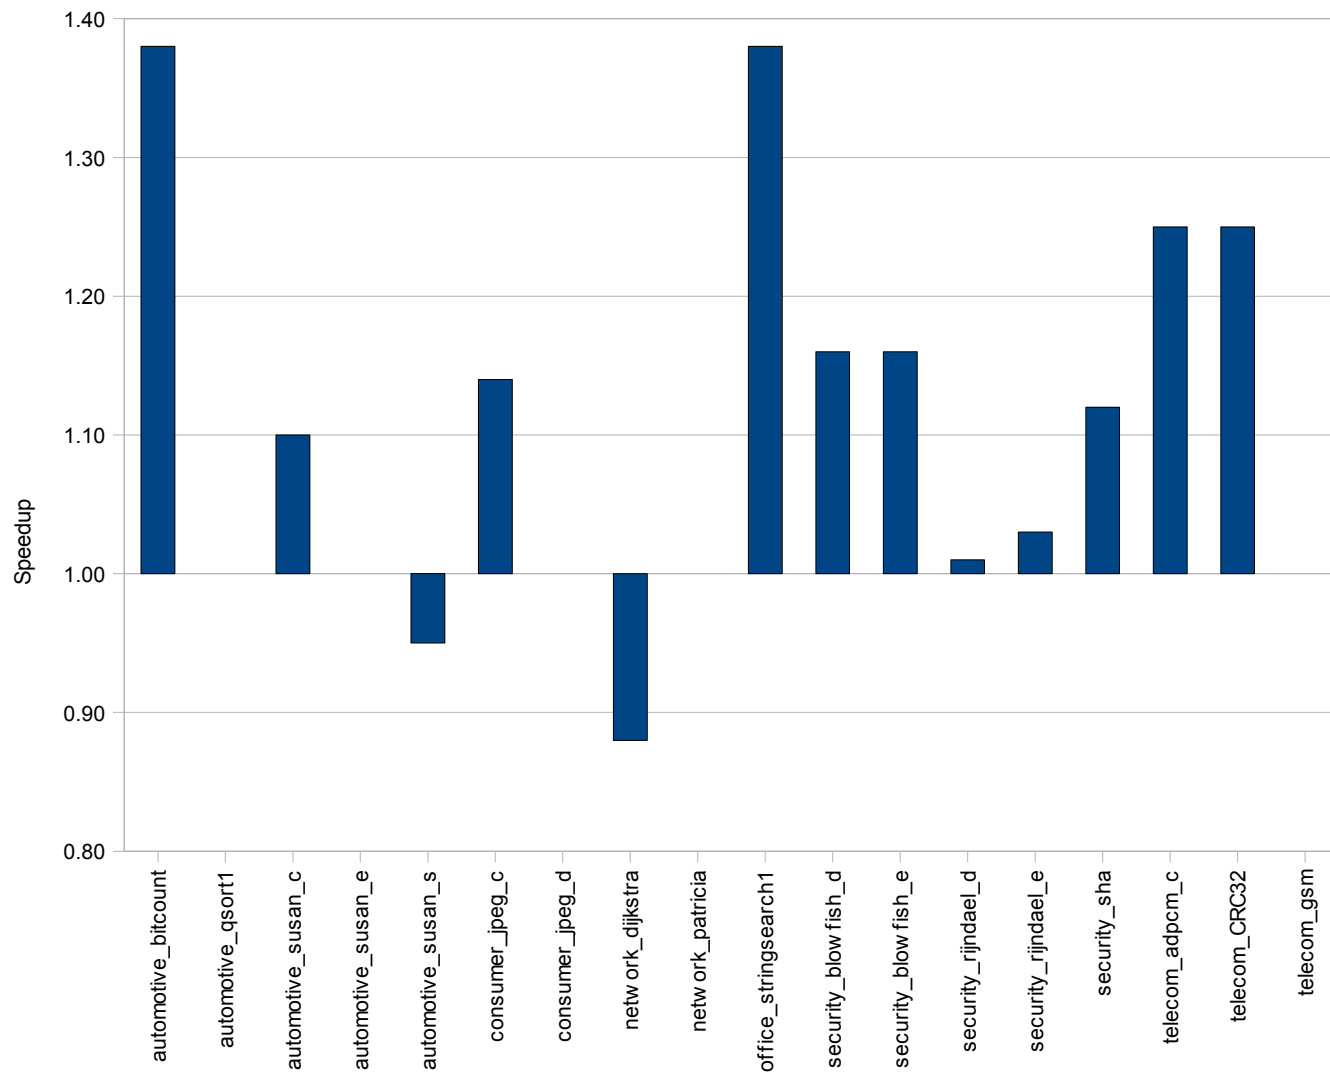
An Experiment Results

Benchmark

Benchmark	Speedup
automotive_bitcount	1.38
automotive_qsort1	1.00
automotive_susan_c	1.10
automotive_susan_e	1.00
automotive_susan_s	0.95
consumer_jpeg_c	1.14
consumer_jpeg_d	1.00
network_dijkstra	0.88
network_patricia	1.00
office_stringsearch1	1.38
security_blowfish_d	1.16
security_blowfish_e	1.16
security_rijndael_d	1.01
security_rijndael_e	1.03
security_sha	1.12
telecom_adpcm_c	1.25
telecom_CRC32	1.25
telecom_gsm	1.00

Average

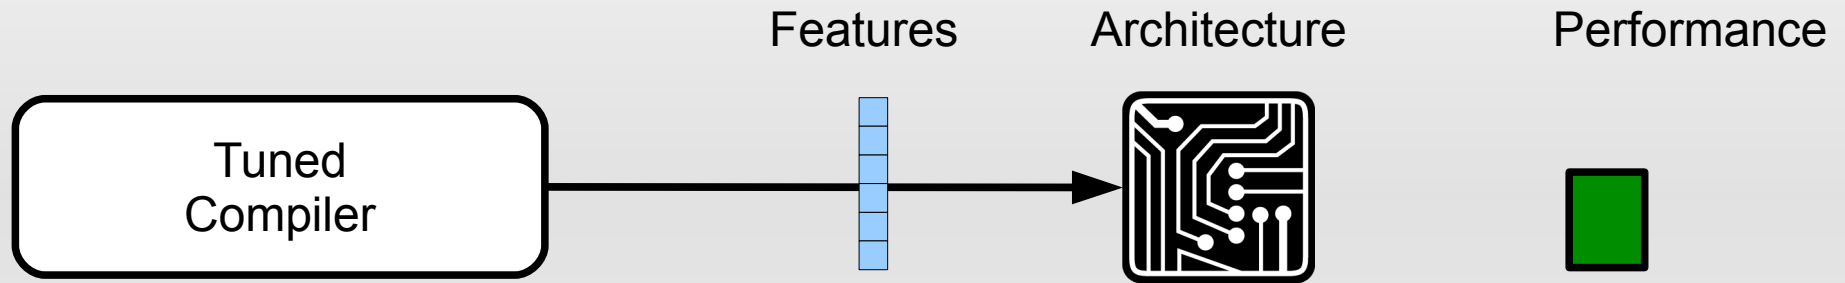
1.10



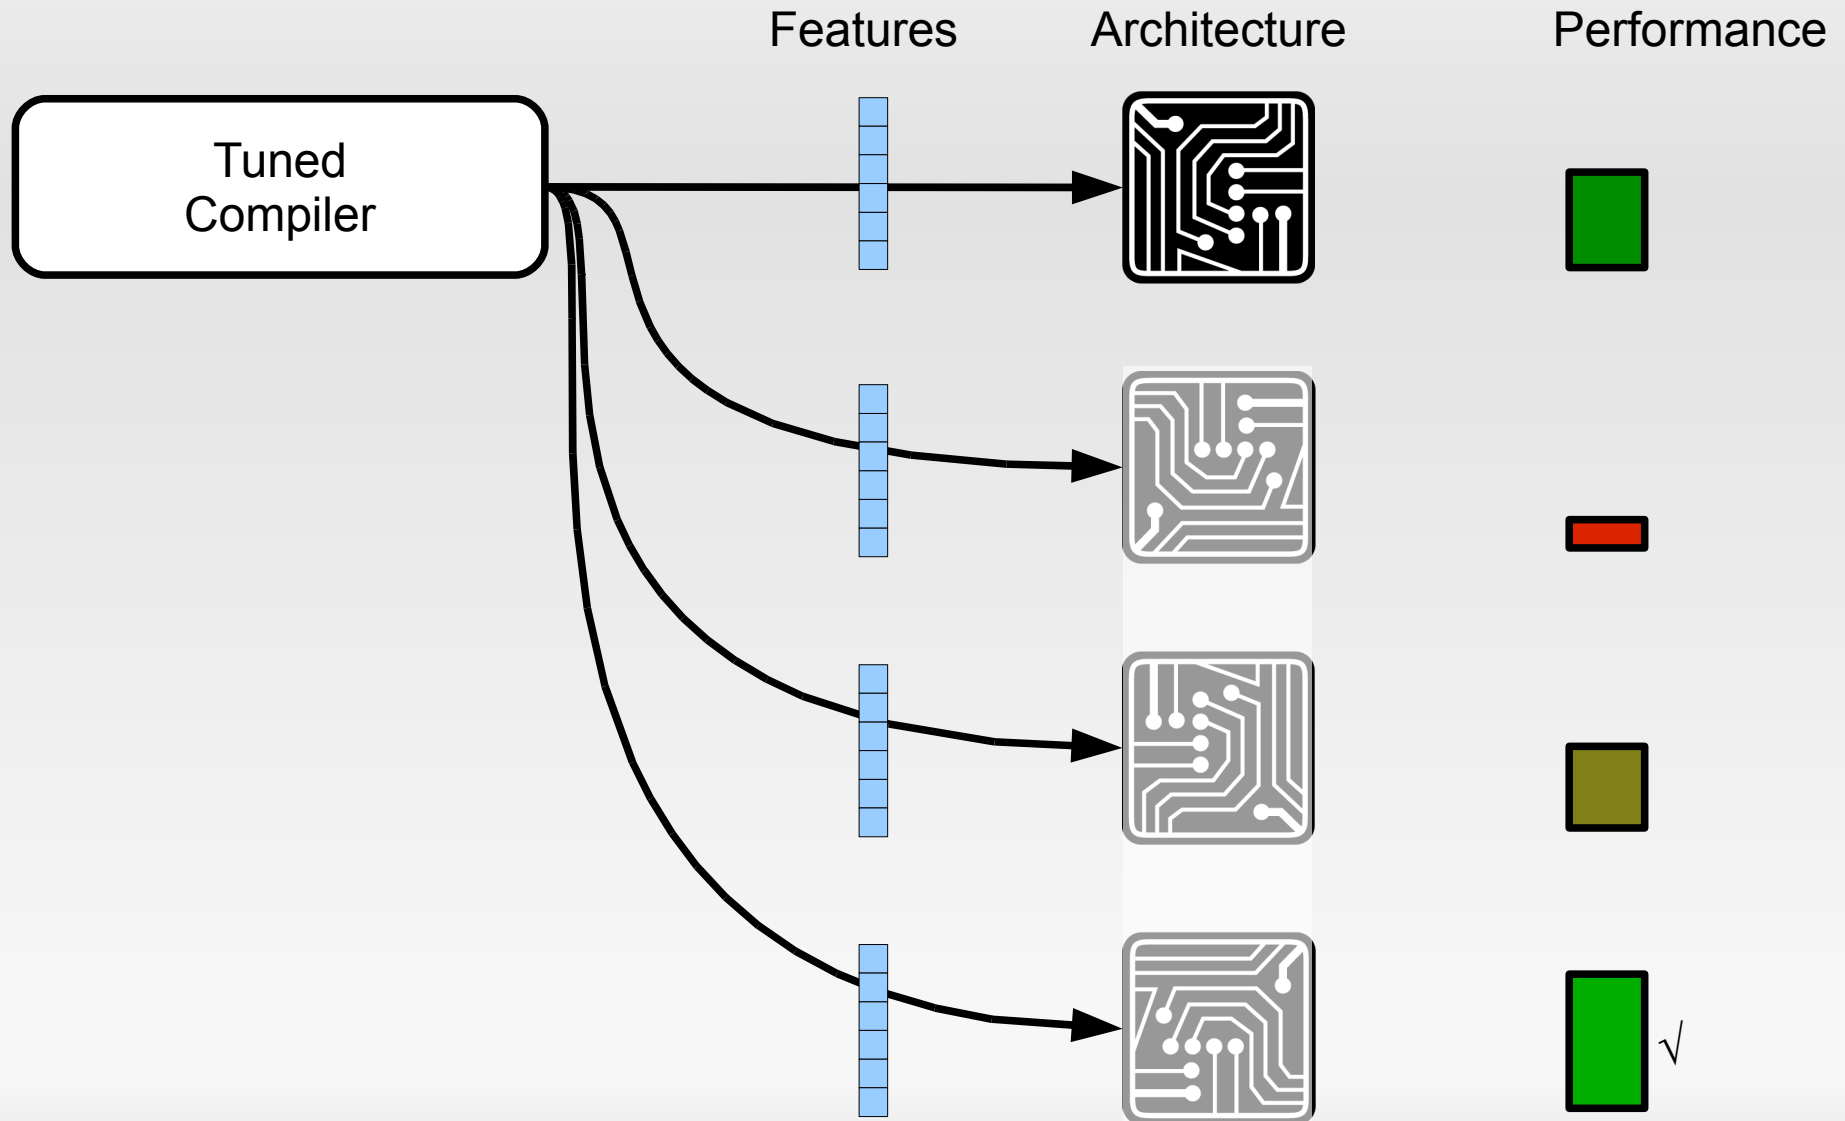
Co-Design Exploration

- Designing an embedded system for some application
- Create many architectures in a 'family'
- Choose the most efficient

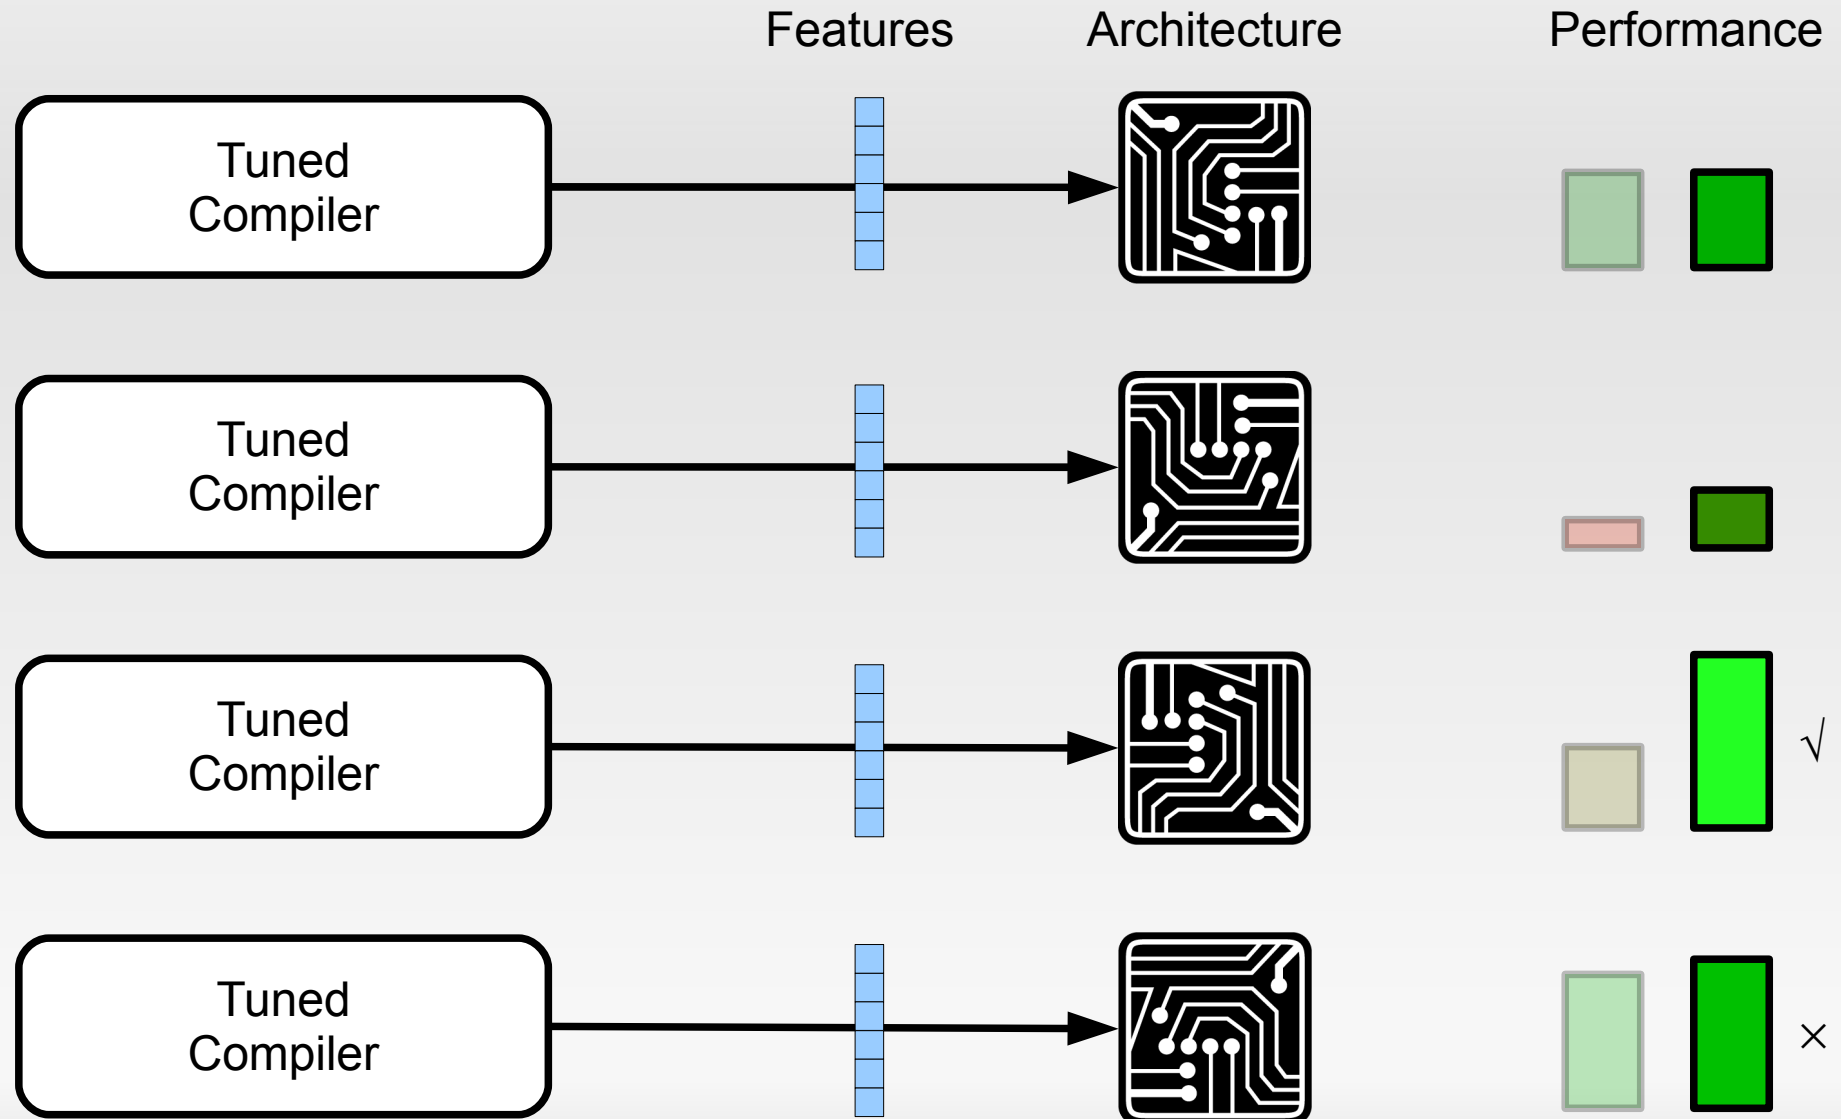
Co-Design Exploration



Co-Design Exploration



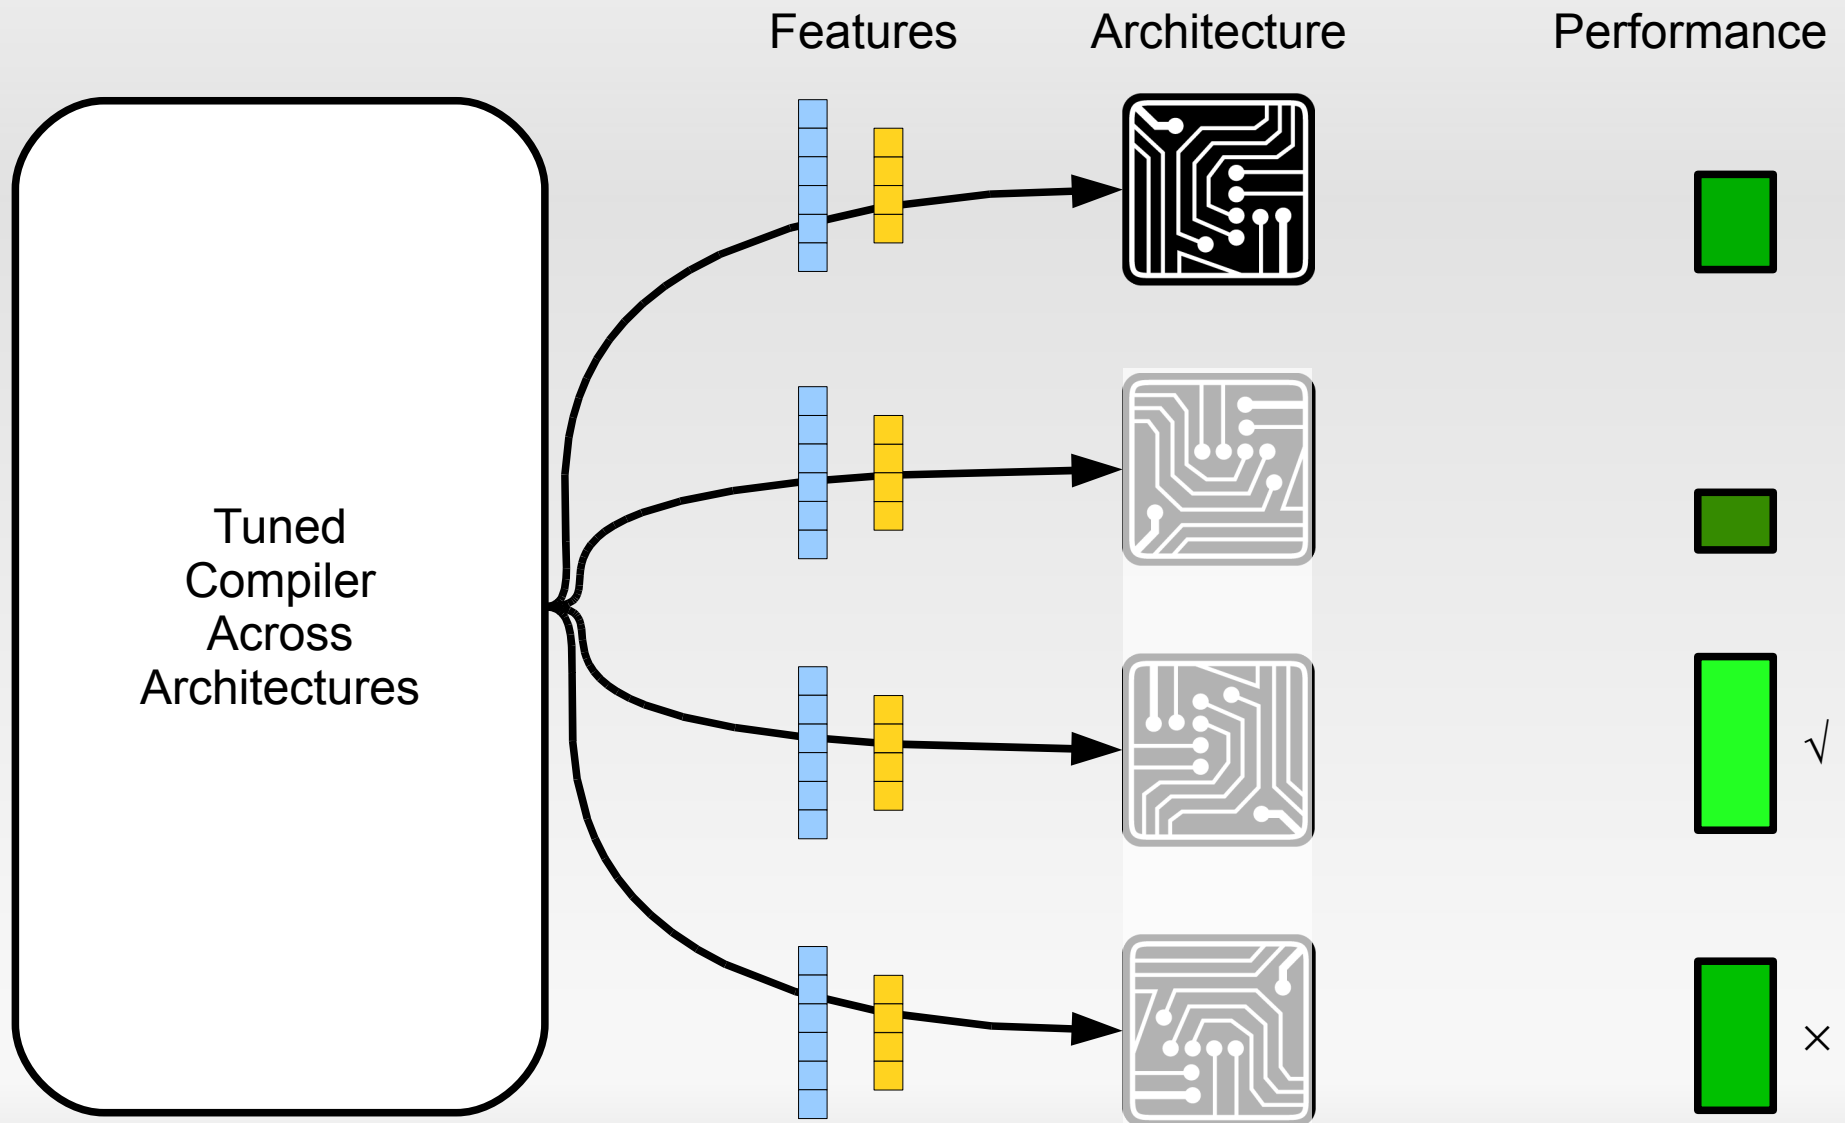
Co-Design Exploration



Co-Design Exploration

- Add some architectural features
 - Cache sizes
 - Number of registers
 - Number & type of FUs
- Incorporate into machine learning for heuristics
- Automatically learned a tuned compiler for a new architecture!

Co-Design Exploration



Open Source Work

- Must have extensible compiler
 - Alter heuristics
 - Gather features
 - Push button compiler tuning
 - Work on configurable processors

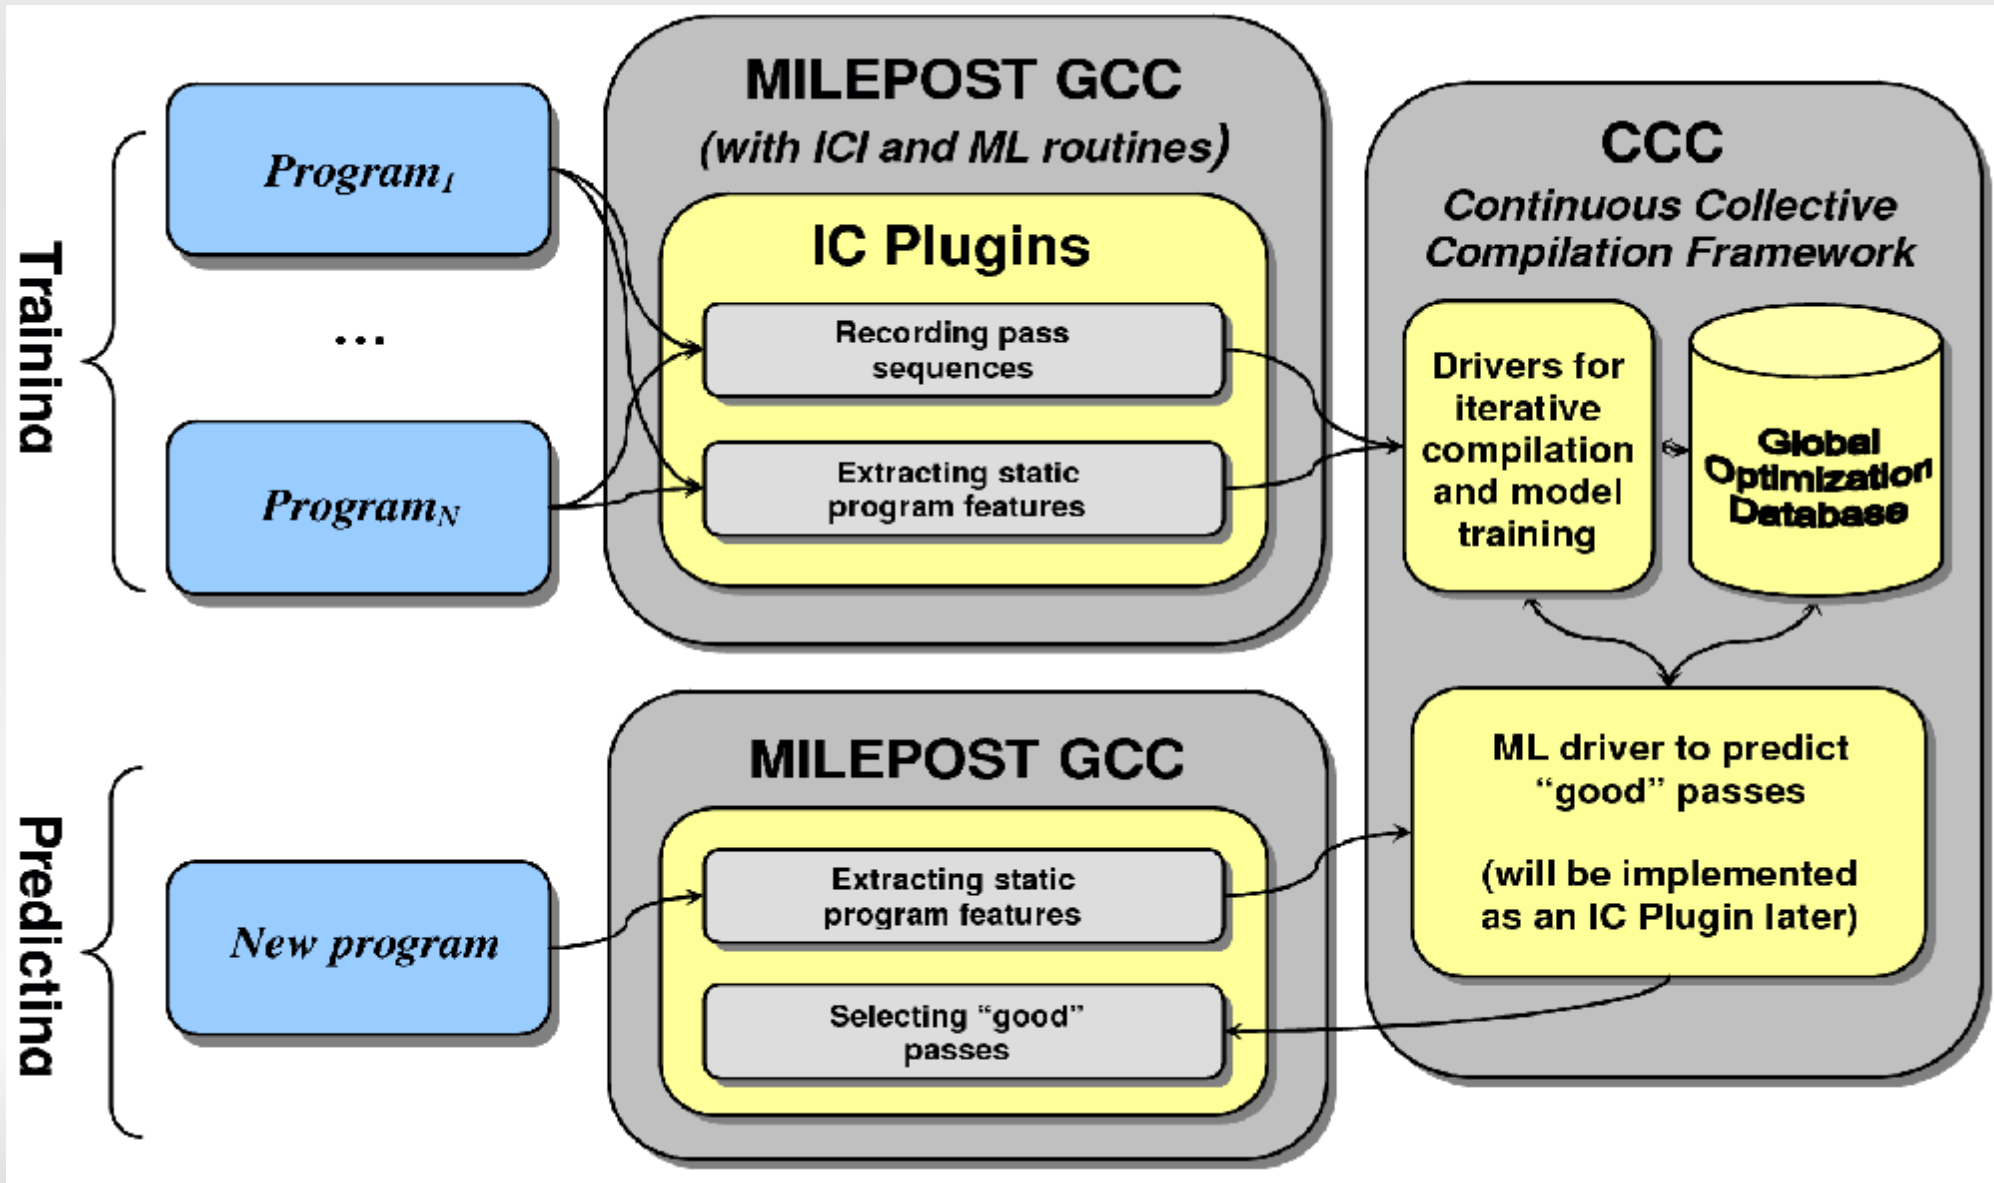
Open Source Work

- MilePost GCC (milepost branch in GCC SVN)
 - Extensibility through shared libraries
 - Extensions to control heuristics
 - Extensions to compute features
 - Provides an 'Iterative Compilation Interface'

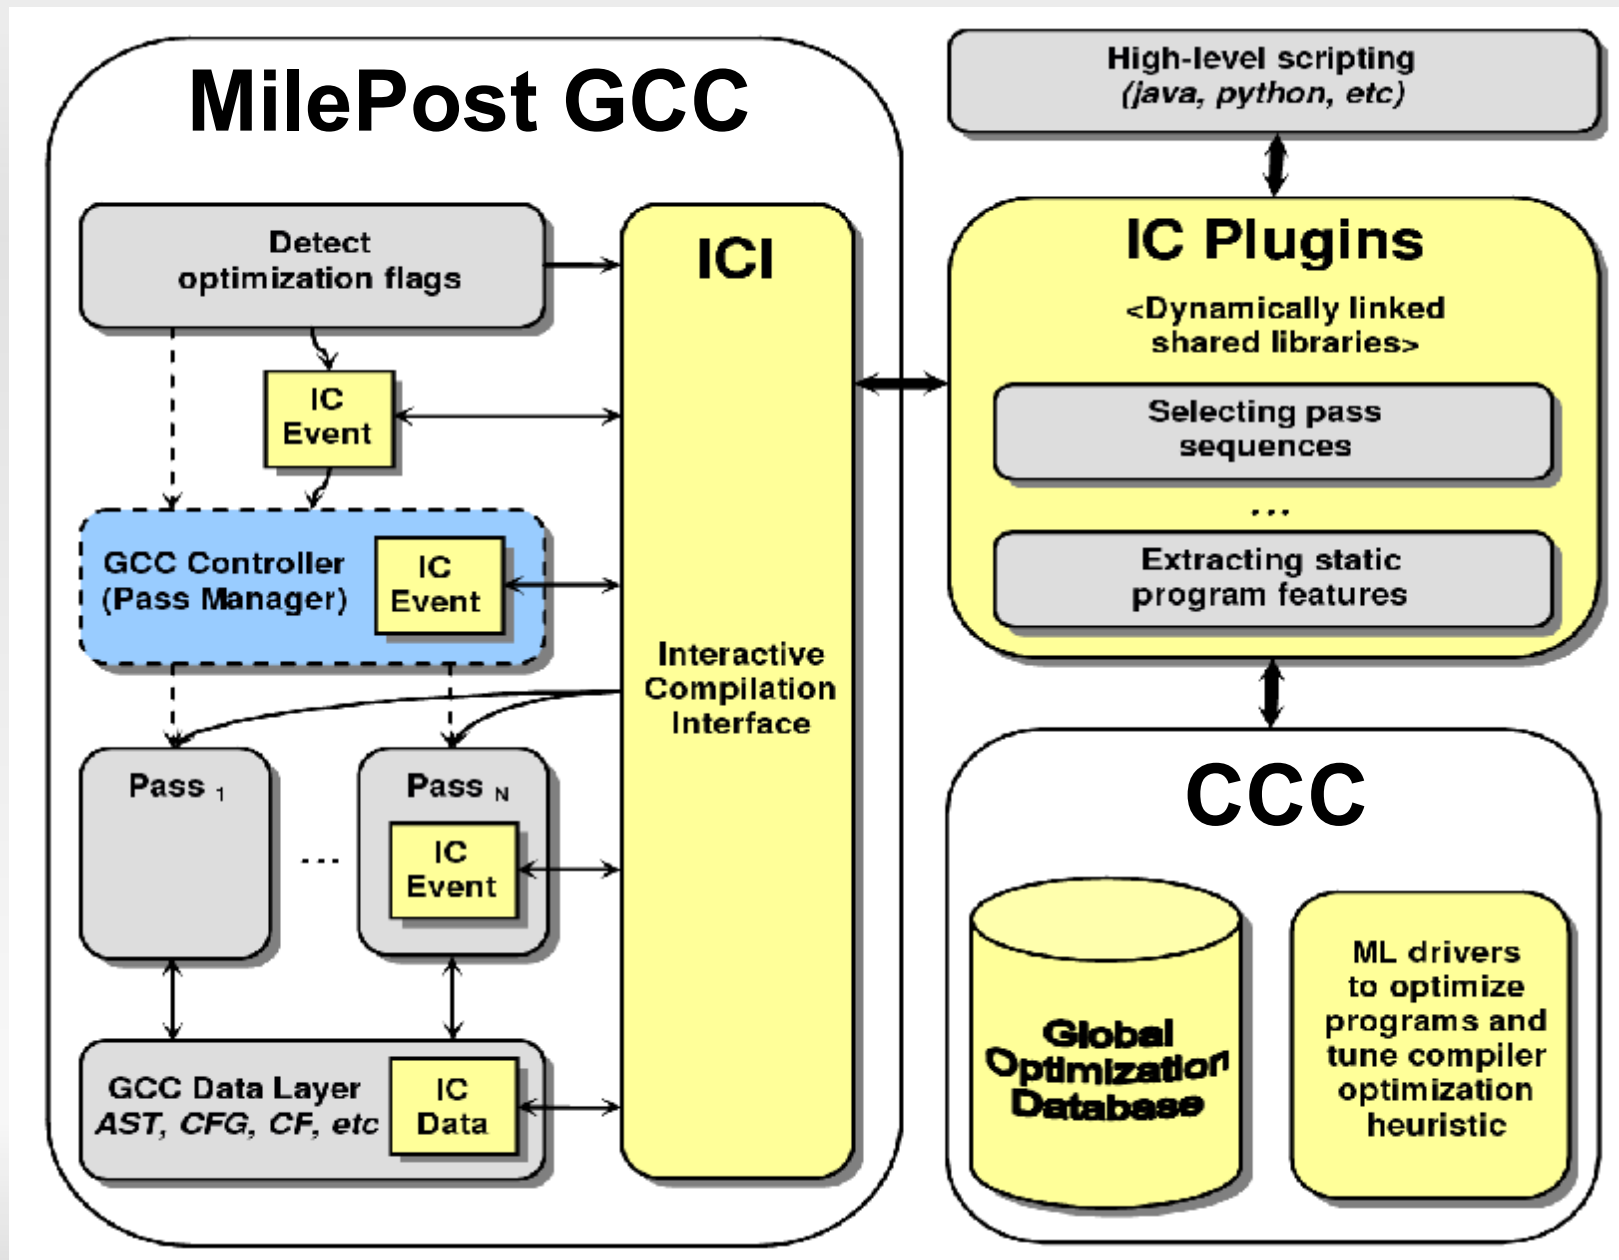
Open Source Work

- Continuous Collective Compilation Framework
(<http://sourceforge.net/projects/cccpf>)
 - Driver framework – automatically tunes compiler
 - Runs iterative compilations
 - Benchmarks and settings included
 - Learns heuristics
 - Heuristics predefined
 - Uses MilePost GCC

Open Source Work



Open Source Work



Open Source Work

- ARC tool chain (<http://www.arc.com/software/development/gnutools.html>)
 - Provides back-end for ARC processors
 - GCC, GDB, BinUtils, uClibc
 - MilePost GCC can be layered on top

OSS Challenges

- Political
 - Extensibility in GCC seen as a threat
 - Major changes need substantial backing
- Legal
 - GPL3 makes lawyers quiver
- How to make money from OSS
- Moving target
 - Just when your code works someone changes the compiler

OSS Benefits

- Huge user base
 - Peer provided support better than most companies'
 - Much wider feedback / testing
 - Allows greater visibility
- Companies can make use of others' work
 - No NDAs needed once OSS approval in house
- Customers expect GCC availability
 - Not so good for compiler vendors

Conclusion

- A system for automatically tuning compilers
- Improves over human created heuristics
- May learn across architectures
- Heavy use of open source software
- Huge and exciting new field opening up

Questions?